

Problem Set 8

Yue Hu

Nov 2017

problem 1

(a)

Parato distribution is a heavy-tailed distribution and its tail decays more slowly to 0.

(b)

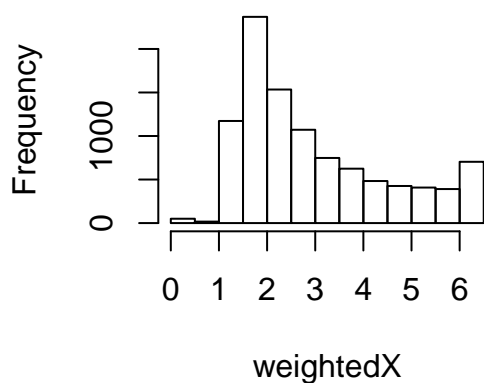
```
library(EnvStats)
# generate sample
sample <- EnvStats::rpareto(10000, 2,3)

##calculate Ex
weightedX <- sapply(sample, function(x) {x*x^4*exp(2-x)/24})
EX <- mean(weightedX)
cat("Ex is", EX)

## Ex is 2.981199


hist(weightedX)
```

Histogram of weightedX



```
#histogram of weight
weight <- sapply(sample, function(x) {x^4*exp(2-x)/24})
hist(weight)
```

Histogram of weight

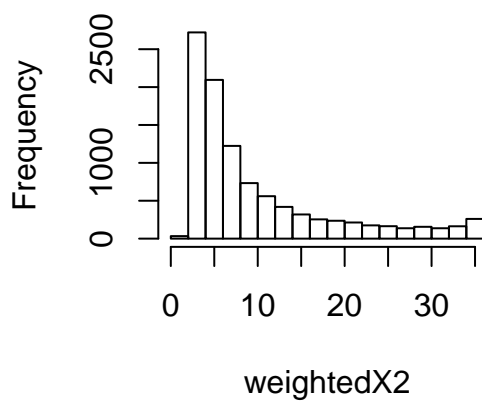


```
# calculate  $E(x^2)$ 
weightedX2 <- sapply(sample, function(x) {x^2*x^4*exp(2-x)/24})
EX2 <- mean(weightedX2)
cat("E(x^2) is", EX2)

## E(x^2) is 9.883008

hist(weightedX2)
```

Histogram of weightedX2



so we can see the value of EX and $E(X^2)$ is as expected.
And the variance is relatively small and resonable.

(c)

```
# generate sample
sample2 <- rexp(10000)
```

```

sample2 <- sample2 +2

#calculate Ex
weightedXp <- sapply(sample2, function(x) {x*24/x^4/exp(2-x)})
EXp <- mean(weightedXp)
cat("EX is ", EXp)

## EX is  2.899011

hist(weightedXp)

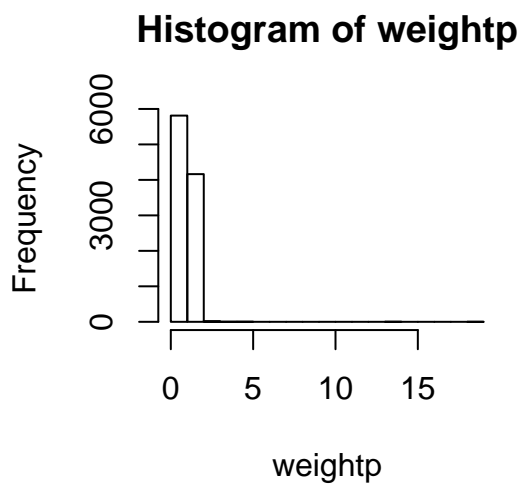
```



```

#histogram of weight
weightp <- sapply(sample2, function(x) {24/x^4/exp(2-x)})
hist(weightp)

```




```

#calculate  $E(X^2)$ 
weightedX2p <- sapply(sample2, function(x) {x^2*24/x^4/exp(2-x)})
EX2p <- mean(weightedX2p)
cat("E(X^2) is ", EX2p)

## E(X^2) is 9.877829

hist(weightedX2p)

```



We can see the variance is much larger than previous case
 IN fact, when g has thinner tails than f , f/g can be unbounded, and the variance of EX can be very large.

problem 2

```


theta <- function(x1,x2) atan2(x2, x1)/(2*pi)

f <- function(x) {
  f1 <- 10*(x[3] - 10*theta(x[1],x[2]))
  f2 <- 10*(sqrt(x[1]^2 + x[2]^2) - 1)
  f3 <- x[3]
  return(f1^2 + f2^2 + f3^2)
}


# plot slices
require(fields)
x1s <- seq(-5, 5, len = 100)
x2s <- seq(-5, 5, len = 100)

#slice with x3 = 10
fx <- apply(cbind(expand.grid(x1s, x2s), 10), 1, f)
image.plot(x1s, x2s, matrix(log(fx), 100, 100))

```



```
# slice with x3 = 1
fx2 <- apply(cbind(expand.grid(x1s, x2s), 1), 1, f)
image.plot(x1s, x2s, matrix(log(fx2), 100, 100))
```



```
# optimize by optim()
init <- c(2,2,2)
optim(init, f, method = 'BFGS', control = list(trace = TRUE))

## initial value 394.564575
## iter 10 value 2.282921
## iter 20 value 0.030005
## final value 0.000000
## converged
## $par
## [1] 1.000000e+00 7.466597e-12 8.048624e-12
##
```

```

## $value
## [1] 4.587856e-20
##
## $counts
## function gradient
##      72      29
##
## $convergence
## [1] 0
##
## $message
## NULL

# use different starting points
init <- c(-200,20000,10000)
optim(init, f, method = 'BFGS', control = list(trace = TRUE))

## initial value 50094968703.060165
## iter 10 value 675.661422
## iter 20 value 0.000653
## final value 0.000000
## converged
## $par
## [1] 1.000000e+00 4.143270e-10 6.504327e-10
##
## $value
## [1] 4.320929e-19
##
## $counts
## function gradient
##      70      29
##
## $convergence
## [1] 0
##
## $message
## NULL

init <- c(2,2, 2)
nlm(f,init)

## $minimum
## [1] 1.70091e-08
##
## $estimate
## [1] 9.999995e-01 -8.223157e-05 -1.300799e-04
##
## $gradient
## [1] -1.343915e-08 1.545268e-08 -9.398712e-09
##
## $code
## [1] 1
##
## $iterations
## [1] 22

```

```

init <- c(2e16, 2e16, -2e16)
nlm(f, init)

## $minimum
## [1] 1.700873e-08
##
## $estimate
## [1] 9.999995e-01 -8.223072e-05 -1.300783e-04
##
## $gradient
## [1] -1.914502e-08 -7.915570e-08 5.333774e-08
##
## $code
## [1] 1
##
## $iterations
## [1] 31

```

so we can see that the optimum parameter for f is (1,0,0). With different starting points they roughly stop at the same point, though some are reaching better minimums.

Problem3

(a)

Set the value of censored data as Z_i ($i=1,2,\dots,c$). And the rest value Y_j ($j=1,2,\dots,n-c$) are known. And the likelihood should be

$$P(Y_i, Z_i | X; \theta) = \prod_i (2\pi\sigma^2)^{-1/2} e^{-\frac{(Z_i - \beta_0 - \beta_1 x_i)^2}{2\sigma^2}} \prod_j (2\pi\sigma^2)^{-1/2} e^{-\frac{(Y_j - \beta_0 - \beta_1 x_j)^2}{2\sigma^2}}$$

and The log-likelihood should be

$$LL(Y_i, Z_i | X; \theta) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_i (Z_i - \beta_0 - \beta_1 x_i)^2 - \frac{1}{2\sigma^2} \sum_j (Y_j - \beta_0 - \beta_1 x_j)^2$$

For the E step, calculate

$$\begin{aligned}
Q &= E(Y_i, Z_i | X; \theta) \\
&= E\left[-\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_i (Z_i - \beta_0 - \beta_1 x_i)^2 - \frac{1}{2\sigma^2} \sum_j (Y_j - \beta_0 - \beta_1 x_j)^2\right] \\
&= -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_j (Y_j - \beta_0 - \beta_1 x_j)^2 - \frac{1}{2\sigma^2} \sum_i E[(Z_i - \beta_0 - \beta_1 x_i)^2] \\
&= -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_j (Y_j - \beta_0 - \beta_1 x_j)^2 - \frac{1}{2\sigma^2} \sum_{i=1}^c [E(Z_i) - \beta_0 - \beta_1 x_i]^2 + \text{Var}(Z_i)] \\
&= -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{j=i}^{n-c} (Y_j - \beta_0 - \beta_1 x_j)^2 - \frac{1}{2\sigma^2} \sum_{i=1}^c (E(Z_i) - \beta_0 - \beta_1 x_i)^2 + c \text{Var}(Z_i)
\end{aligned}$$

where $E(Z_i)$ and $\text{Var}(Z_i)$ is the expectation and variance of Z_i based on the current value at iteration t , and should be constant.

For M step, optimization is with respect to β_0, β_1, σ .

This is similar to linear regression and result should be

$$\hat{\beta}_0 = \frac{\sum_i (x_i - \bar{x})(Y_i - \bar{Y})}{\sum_i (x_i - \bar{x})^2}$$

$$\hat{\beta}_1 = \bar{Y} - \hat{\beta}_0 \bar{x}$$

$$\hat{\sigma}^2 = \frac{\sum_{j=i}^{n-c} (Y_i - \hat{\beta}_0 - \hat{\beta}_1 x_j)^2 + c(E_z - \hat{\beta}_0 - \hat{\beta}_1 x_j)^2 + cVar(Z)}{n}$$

Where Y_i is the uncensored data plus c censored data with values $E(z)$

(b)

For initial value we can we can ignore datas where $y > \tau$, and use lm on the truncated y sequence.

(c)

For the censored data we can use $E(W|W > \tau)$ for y values, then use lm on the new y sequence to get betas. For sigma, the first two terms is the σ^2 by lm, and we need only add $V(z)/n$ to it and take the square root.

We should break the optimization if objective function ceases to change. In this case, when the previous state and the current state considered by the EM algorithm are not significantly large.

```
# simulate data
set.seed(1)
n <- 100
beta0 <- 1
beta1 <- 2
sigma2 <- 6

x <- runif(n)
yComplete <- rnorm(n, beta0 + beta1*x, sqrt(sigma2))

# censor the sequence, 20% exceedance
tau1 <- quantile(yComplete, .8)
c <- 20

# calculate initial value

yTruncate <- yComplete
yTruncate[yTruncate>tau1] <- NA
xTruncate <- x
xTruncate[yTruncate>tau1] <- NA
mod <- lm(yTruncate~xTruncate)
betai <- summary(mod)$coef
sigmai <- summary(mod)$sigma

##
step = 0
for (i in 1:1000){
  # calculate Expectation and variance of Z based on current theta
  ymean <- betai[1]+betai[2]*mean(x)
  tauStar <- (tau1-ymean)/sigmai
  thou <- dnorm(tauStar)/(1-pnorm(tauStar))
  EZ <- ymean + sigmai* thou
```



```

varZ <- sigmai^2*(1+tauStar*thou-thou^2)

# construct new y sequence
yNew <- yComplete
yNew[yNew>tau1] <- EZ
mod <- lm(yNew~x)
# break the optimization if the difference between new and old parameters are small enough
if (abs(summary(mod)$coef[1]-betai[1]) < 1e-16 & abs(summary(mod)$coef[2]-betai[2])< 1e-16)
{break}
# update theta
betai <- summary(mod)$coef
sigmai <-sqrt(summary(mod)$sigma + c*varZ/n)
step <- step + 1
}

cat("theta is {" ,c(betai[1],betai[2], sigmai), "}, taking", step, "steps")

## theta is { 0.4855346 2.534128 1.421644 }, taking 17 steps

```

for sequence of 80% exceedance

```

set.seed(1)
n <- 100
beta0 <- 1
beta1 <- 2
sigma2 <- 6

x <- runif(n)
yComplete <- rnorm(n, beta0 + beta1*x, sqrt(sigma2))

# censor the sequence, 80%exceedance
tau2 <- quantile(yComplete, 0.2)
c <- 80
# calculate initial value

yTruncate <- yComplete
yTruncate[yTruncate>tau2] <- NA
xTruncate <- x
xTruncate[yTruncate>tau2] <- NA
mod <- lm(yTruncate~xTruncate)
betai <- summary(mod)$coef
sigmai <- summary(mod)$sigma

##
step = 0
for (i in 1:1000){
  yEstimate <- sapply(x, function(x) betai[1]+betai[2]*x)
  mean <- mean(yEstimate)
  tauStar <- (tau2-mean)/sigmai
  thou <- dnorm(tauStar)/(1-pnorm(tauStar))
  EZ <- mean + sigmai* thou
  varZ <- sigmai^2*(1+tauStar*thou-thou^2)

```

```

# construct new y sequence
yNew <- yComplete
yNew[yNew>tau2] <- EZ
mod <- lm(yNew~x)
# break the optimization if the difference between new and old parameters are small enough
if (abs(summary(mod)$coef[1]-betai[1]) < 1e-16 & abs(summary(mod)$coef[2]-betai[2])< 1e-16)
{break}
betai <- summary(mod)$coef
sigmai <-sqrt(summary(mod)$sigma + c*varZ/n)
step <- step + 1

}

cat("theta is {" ,c(betai[1],betai[2], sigmai), "}, taking", step, "steps")

## theta is { 0.4725494 1.253678 1.458363 }, taking 131 steps

```

(d)

```

# simulate data
set.seed(1)
n <- 100
beta0 <- 1
beta1 <- 2
sigma2 <- 6

x <- runif(n)
yComplete <- rnorm(n, beta0 + beta1*x, sqrt(sigma2))

# generate truncated sequence, 20%exceedance
tau1 <- quantile(yComplete, .8)

# construct truncated sequece
yTruncate <- yComplete
xTruncate <- x
xTruncate <- xTruncate[yTruncate < tau1]
yTruncate <- yTruncate[yTruncate < tau1]
n<- length(xTruncate)

# theta being {beta0, beta1, log(sigma) }
LL <- function(theta, x, y){
  yEstimate <- sapply (x, function(x) theta[1]+theta[2]*x)
  rss <-sum((yEstimate-y)^2)
  return(-n/2*log(2*3.14)-n*theta[3]-0.5*exp(-2*theta[3])*rss)
}

optim(c(0,0,0), LL, y = yTruncate,x = xTruncate, method = 'BFGS', control = list(trace = TRUE, parscale=1000))

## initial value -233.200173
## final value -23770491988660368346608028266066620642448082888048282048642088220824408460020422426240

```

```

## converged
## $par
## [1] -206.6395 -116.0925 -185.2685
##
## $value
## [1] -2.377049e+167
##
## $counts
## function gradient
##      376      4
##
## $convergence
## [1] 0
##
## $message
## NULL

# try several initial points to avoid local minima
optim(c(1.5,1.5,10), LL, y = yTruncate,x = xTruncate, method = 'BFGS', control = list(trace = TRUE, parsca

## initial value -873.494800
## final value -118842243768032342058428.000000
## converged
## $par
## [1] 2.814750e+11 2.814750e+11 1.485528e+21
##
## $value
## [1] -1.188422e+23
##
## $counts
## function gradient
##      7      7
##
## $convergence
## [1] 0
##
## $message
## NULL

optim(c(1,2,0.4), LL, y = yTruncate,x = xTruncate, method = 'BFGS', control = list(trace = TRUE, parsca

## initial value -165.483237
## final value -69539519714375289894482086484660266068226846880220882026484062442006422644688884488446
## converged
## $par
## [1] 684.7854 356.0316 -268.6741
##
## $value
## [1] -6.953952e+240
##
## $counts
## function gradient
##      366      3
##
## $convergence

```

```
## [1] 0
##
## $message
## NULL

optim(c(1,2,10), LL, y = yTruncate,x = xTruncate, method = 'BFGS', control = list(trace = TRUE, parscale = 1e-05))

## initial value -873.494799
## final value -873.494799
## converged
## $par
## [1] 1 2 10
##
## $value
## [1] -873.4948
##
## $counts
## function gradient
##      2      1
##
## $convergence
## [1] 0
##
## $message
## NULL
```

So we can see that with $\theta = 1, 2, 10$, the log-likelihood reaches its maximum.
And EM algorithm uses less steps.